

User Manual For Real-Time Sensor Data Streaming

Harmit Khimani & Nishank Kansara

B.Tech. Mini Project - Under Prof. Tapas Kumar Maiti

Contents

1	Prerequisites	2
2	Hardware Setup & Wiring	2
3	ESP32 MQTT Publisher Code	3
3.1	Install Libraries	3
3.2	Embed MQTT Credentials	3
3.3	Obtain HiveMQ Cloud Details	3
4	Kafka Setup	4
4.1	Disable Firewall	4
4.2	Download Kafka 3.9.0	4
4.3	Configure Kafka Listeners	4
4.4	Start ZooKeeper	4
4.5	Start Kafka Server	4
4.6	Create Verify Topic	4
5	Data Storage, Visualization & Smart Feedback	6
5.1	PostgreSQL (PGAdmin)	6
5.2	Web Visualization	7
5.3	Real-Time Smart Feedback	8
6	Additional Resources	9
7	Troubleshooting	9

1 Prerequisites

Ensure you have the following installed and configured:

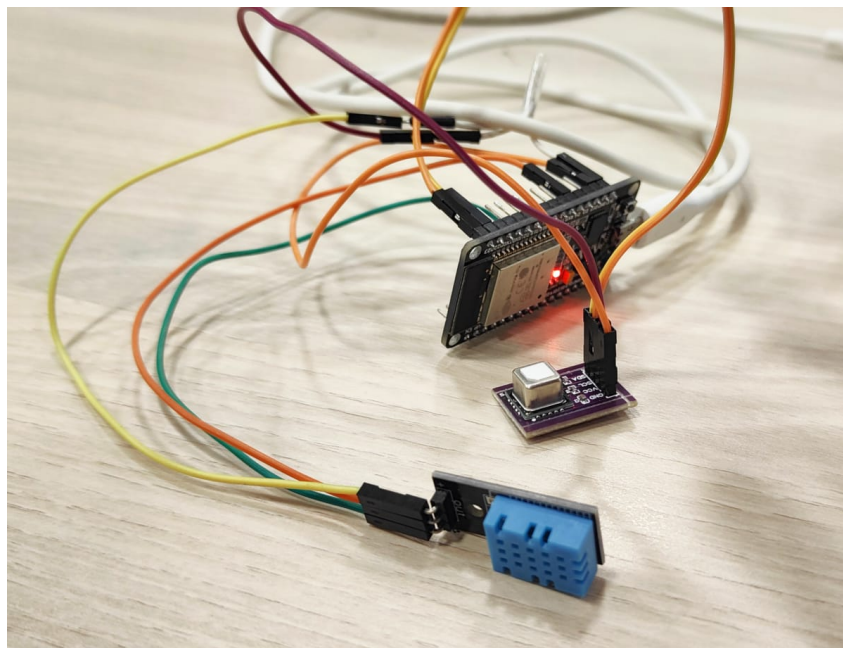
- **Java 11+**
- **Apache Kafka 3.9.0 (Scala 2.13)**
- **ESP32 toolchain** (Arduino IDE or PlatformIO)
- Network connectivity: ESP32, Kafka server, and Spring Boot host on the same LAN or Wi-Fi.

2 Hardware Setup & Wiring

Connect sensors directly to the ESP32 via USB power.

Sensor	ESP32 Pin	Notes
DHT11	GPIO 4	Data line (10 K Ω pull-up)
SCD40	SDA \rightarrow GPIO 22, SCL \rightarrow GPIO 23	I ² C bus lines
VCC	3V3	3.3 V from ESP32
GND	GND	Common ground

Table 1: Sensor to ESP32 Wiring



3 ESP32 MQTT Publisher Code

3.1 Install Libraries

Add these libraries to your Arduino IDE:

- **PubSubClient** for MQTT
- **DHT** sensor library
- **SCD4x** library for CO₂ sensing
- **Wire** for I²C communication

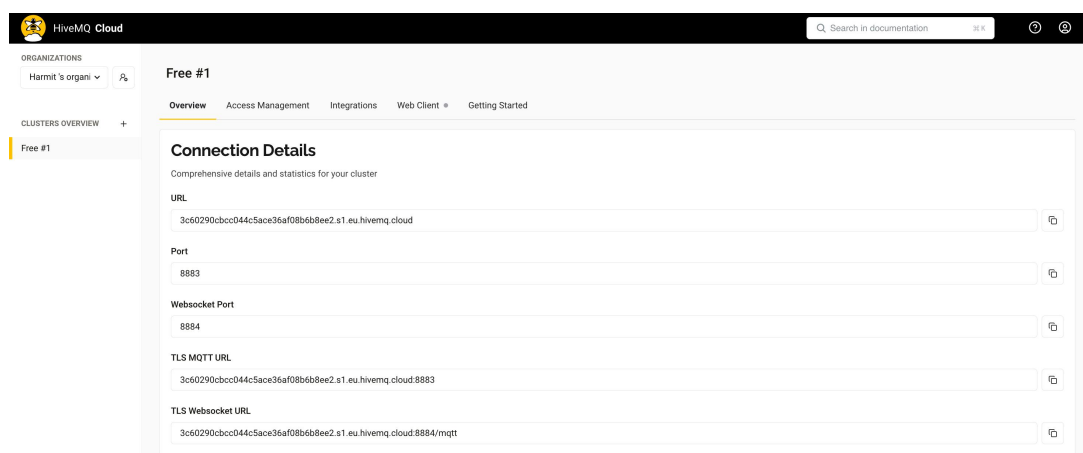
3.2 Embed MQTT Credentials

```
// W i F i
const char* ssid      = "<YOUR_SSID>";
const char* password  = "<YOUR_WIFI_PASSWORD>";

// HiveMQ Cloud
const char* mqtt_server = "<YOUR_MQTT_SERVER_URL>";
const int   mqtt_port   = 8883;
const char* mqtt_username = "<USERNAME>";
const char* mqtt_password = "<PASSWORD>";
const char* topic       = "<TOPIC>";
```

3.3 Obtain HiveMQ Cloud Details

1. Sign in at <https://console.hivemq.cloud>.
2. Create a free cluster.
3. Copy the **Cluster Endpoint** and paste it into the `mqtt_server` configuration field.
4. Retrieve or generate your **Credentials** for a secure connection.



4 Kafka Setup

4.1 Disable Firewall

Open port 9092:

Windows: netsh advfirewall firewall add rule name="KafkaPort" dir=in action=allow protocol=TCP localport=9092

macOS (Terminal): sudo /usr/libexec/ApplicationFirewall/socketfilterfw --setglobalstate off

4.2 Download Kafka 3.9.0

1. Visit <https://kafka.apache.org/downloads>.
2. Select Scala 2.13 + Kafka 3.9.0.
3. Extract the archive.

4.3 Configure Kafka Listeners

In config/server.properties:

```
listeners=PLAINTEXT://0.0.0.0:9092
advertised.listeners=PLAINTEXT://<YOUR_HOST_IP>:9092
```

4.4 Start ZooKeeper

```
cd kafka_2.13-3.9.0
bin/zookeeper-server-start.sh config/zookeeper.properties
```

4.5 Start Kafka Server

```
cd kafka_2.13-3.9.0
bin/kafka-server-start.sh config/server.properties
```

4.6 Create Verify Topic

```
# Create topic
bin/kafka-topics.sh --create \
  --topic esp32-mqtt-topic \
  --bootstrap-server <BROKER_IP>:9092 \
  --partitions 1 --replication-factor 3

# List topics
echo "Available topics:" && \
bin/kafka-topics.sh --list --bootstrap-server localhost:9092

# Test consumption
bin/kafka-console-consumer.sh \
  --bootstrap-server <BROKER_IP>:9092 \
  --topic esp32-mqtt-topic \
  --from-beginning
```

```
harmit ~ - ssh - 80x26
last login: Sun May 4 14:53:53 on tty003
harmit@Mac ~ % sudo /usr/libexec/ApplicationFirewall/socketfilterfw --setglobalstate off
Password:
harmit@Mac ~ %

kafka_2.13-3.0.0 -- java -Xmx1G -Xms1G -server -XX:+UseG1GC -XX:Max...
[2025-05-04 17:54:36,937] INFO Loaded member MemberMetadata(memberId=consumer-mq
tt-data-group-1-b9797d83-633b-4129-989c-c6b9943dc886, groupInstanceId=None, clie
ntId=consumer-mqtt-data-group-1, clientHost/192.168.144.42, sessionTimeoutMs=45
000, rebalanceTimeoutMs=300000, supportedProtocols=List(range)) in group mqtt-da
ta-group with generation 192. (kafka.coordinator.group.GroupMetadata$)
[2025-05-04 17:54:36,937] INFO Loaded member MemberMetadata(memberId=consumer-mq
tt-data-group-1-a8373e79-81f8-4fd0-9998-bfa692ec937b, groupInstanceId=None, clie
ntId=consumer-mqtt-data-group-1, clientHost/192.168.144.42, sessionTimeoutMs=45
000, rebalanceTimeoutMs=300000, supportedProtocols=List(range)) in group mqtt-da
ta-group with generation 192. (kafka.coordinator.group.GroupMetadata$)
[2025-05-04 17:54:36,937] INFO Loaded member MemberMetadata(memberId=consumer-mq
tt-data-group-1-a8373e79-81f8-4fd0-9998-bfa692ec937b, groupInstanceId=None, clie
ntId=consumer-mqtt-data-group-1, clientHost/192.168.144.42, sessionTimeoutMs=45
000, rebalanceTimeoutMs=300000, supportedProtocols=List(range)) in group mqtt-da
ta-group with generation 193. (kafka.coordinator.group.GroupMetadata$)
[2025-05-04 17:54:36,937] INFO [GroupCoordinator 0]: Loading group metadata for
mqtt-data-group with generation 194 (kafka.coordinator.group.GroupCoordinator)
[2025-05-04 17:54:36,937] INFO [GroupMetadataManager brokerId=0] Finished loadin
g offsets and group metadata from _consumer_offsets-13 in 28 milliseconds for e
ach 0, of which 17 milliseconds was spent in the scheduler. (kafka.coordinator
.group.GroupMetadataManager)
[2025-05-04 17:54:36,938] INFO [GroupMetadataManager brokerId=0] Finished loadin
g offsets and group metadata from _consumer_offsets-28 in 29 milliseconds for e
ach 0, of which 29 milliseconds was spent in the scheduler. (kafka.coordinator
.group.GroupMetadataManager)

kafka_2.13-3.0.0 -- java -Xmx512M -Xms512M -server -XX:+UseG1GC -X...
[2025-05-04 17:54:28,190] INFO zookeeper.snapshot.compression.method = CHECKED (
org.apache.zookeeper.server.persistence.SnapStream)
[2025-05-04 17:54:28,193] INFO Reading snapshot /tmp/zookeeper/version-2/snapsho
t.1ae (org.apache.zookeeper.server.persistence.FileSnap)
[2025-05-04 17:54:28,198] INFO The digest in the snapshot has digest version of
2, with zxid as 0x1c2, and digest value as 302547319185 (org.apache.zookeeper.se
rver.DataTree)
[2025-05-04 17:54:28,199] INFO ZooKeeper audit is disabled. (org.apache.zookeep
er.audit.ZKAuditProvider)
[2025-05-04 17:54:28,199] INFO 20 txns loaded in 5 ms (org.apache.zookeeper.serv
er.persistence.FileTxnSnapLog)
[2025-05-04 17:54:28,199] INFO Snapshot loaded in 21 ms, highest zxid is 0x1c2,
digest is 305765972428 (org.apache.zookeeper.server.ZKDatabase)
[2025-05-04 17:54:28,199] INFO Snapshotting: 0x1c2 to /tmp/zookeeper/version-2/s
napshot.1c2 (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2025-05-04 17:54:28,151] INFO Snapshot taken in 2 ms (org.apache.zookeeper.serv
er.ZooKeeperService)
[2025-05-04 17:54:28,155] INFO PrepRequestProcessor (sid:0) started, reconfigEna
bled=false (org.apache.zookeeper.server.PreRequestProcessor)
[2025-05-04 17:54:28,155] INFO zookeeper.request.throttler.shutdownTimeout = 180
00 ms (org.apache.zookeeper.server.RequestThrottler)
[2025-05-04 17:54:28,161] INFO Using checkIntervalMs=60000 maxPerMinute=10000 ma
xNeverUsedIntervalMs=0 (org.apache.zookeeper.server.ContainerManager)
[2025-05-04 17:54:35,846] INFO Creating new log file: log-1c3 (org.apache.zooke
per.server.persistence.FileTxnLog)

harmit@Mac ~ % cd ~/Downloads/kafka_2.13-3.0.0
harmit@Mac ~ % bin/kafka-topics.sh \
--create \
--topic mqtt_data_topic \
--bootstrap-server 192.168.144.205:9092 \
--partitions 1 \
--replication-factor 3
WARNING: Due to limitations in metric names, topics with a period ('.') or under
score ('_') could collide. To avoid issues it is best to use either, but not bot
h.
Error while executing topic command : Topic 'mqtt_data_topic' already exists.
[2025-05-04 17:55:06,445] ERROR org.apache.kafka.common.errors.TopicExistsExcept
(org.apache.kafka.tools.TopicCommand)
harmit@Mac ~ % kafka_2.13-3.0.0 %
```

```
RT RealTimeSensorDataStream Version control
Project
  RealTimeSensorDataStream ~/Desktop/RealTimeSensorDataStream
    idea
    KafkaConsumer
    KafkaProducer
      GetMQTT-SendKafka (miniProject)
        idea
        mvn
        src
        main
        java
          kafka.producer.GetMQTT_SendKafka

Services
  Spring Boot
    GetMQTTSendKafkaApplication @2020/
    MakeDecisionApplication
    StoreDataApplication
    VisualiseDataApplication [devtools]
    Docker

Console
transaction.timeout.ms = 60000
transactional.id = null
value.serializer = class org.apache.kafka.common.serialization.StringSerializer

2025-05-04T17:57:14.497+05:30 INFO 27523 --- [GetMQTT-SendKafka] [ing-boot-client] o.a.k.c.t.i.KafkaMetricsFollow
2025-05-04T17:57:14.583+05:30 INFO 27523 --- [GetMQTT-SendKafka] [ing-boot-client] o.a.k.c.t.i.KafkaMetricsFollow
2025-05-04T17:57:14.521+05:30 INFO 27523 --- [GetMQTT-SendKafka] [ing-boot-client] o.a.k.c.common.utils.AppInfo
2025-05-04T17:57:14.522+05:30 INFO 27523 --- [GetMQTT-SendKafka] [ing-boot-client] o.a.k.c.common.utils.AppInfo
2025-05-04T17:57:14.522+05:30 INFO 27523 --- [GetMQTT-SendKafka] [ing-boot-client] o.a.k.c.common.utils.AppInfo
2025-05-04T17:57:14.653+05:30 INFO 27523 --- [GetMQTT-SendKafka] [afka-producer-1] org.apache.kafka.clients.Metak
Published to Kafka: {"scd_co2":1029,"scd_temp":23.5,"scd_humidity":52.8,"dht_temp":24.1,"dht_humidity":47.0}
2025-05-04T17:57:14.766+05:30 INFO 27523 --- [GetMQTT-SendKafka] [afka-producer-1] o.a.k.c.p.internals.Transaction
Received message: {"scd_co2":1045,"scd_temp":23.5,"scd_humidity":53.2,"dht_temp":24.1,"dht_humidity":47.0}
Published to Kafka: {"scd_co2":1045,"scd_temp":23.5,"scd_humidity":53.2,"dht_temp":24.1,"dht_humidity":47.0}
Received message: {"scd_co2":1050,"scd_temp":23.4,"scd_humidity":53.5,"dht_temp":24.1,"dht_humidity":47.0}
Published to Kafka: {"scd_co2":1050,"scd_temp":23.4,"scd_humidity":53.5,"dht_temp":24.1,"dht_humidity":47.0}
Received message: {"scd_co2":1050,"scd_temp":23.3,"scd_humidity":53.8,"dht_temp":24.1,"dht_humidity":47.0}
Published to Kafka: {"scd_co2":1052,"scd_temp":23.3,"scd_humidity":53.8,"dht_temp":24.1,"dht_humidity":47.0}
Received message: {"scd_co2":1053,"scd_temp":23.3,"scd_humidity":54.2,"dht_temp":24.1,"dht_humidity":47.0}
Published to Kafka: {"scd_co2":1053,"scd_temp":23.3,"scd_humidity":54.2,"dht_temp":24.1,"dht_humidity":47.0}
Received message: {"scd_co2":1053,"scd_temp":23.3,"scd_humidity":54.2,"dht_temp":24.1,"dht_humidity":47.0}
Published to Kafka: {"scd_co2":1053,"scd_temp":23.3,"scd_humidity":54.2,"dht_temp":24.1,"dht_humidity":47.0}

RealTimeSensorDataStream > KafkaProducer > GetMQTT-SendKafka > src > main > java > kafka > producer > GetMQTT_SendKafka > GetMQTTSendKafkaApplication 4/8 LF UTF-8 Tab
```

5 Data Storage, Visualization & Smart Feedback

5.1 PostgreSQL (PGAdmin)

Store historical sensor data:

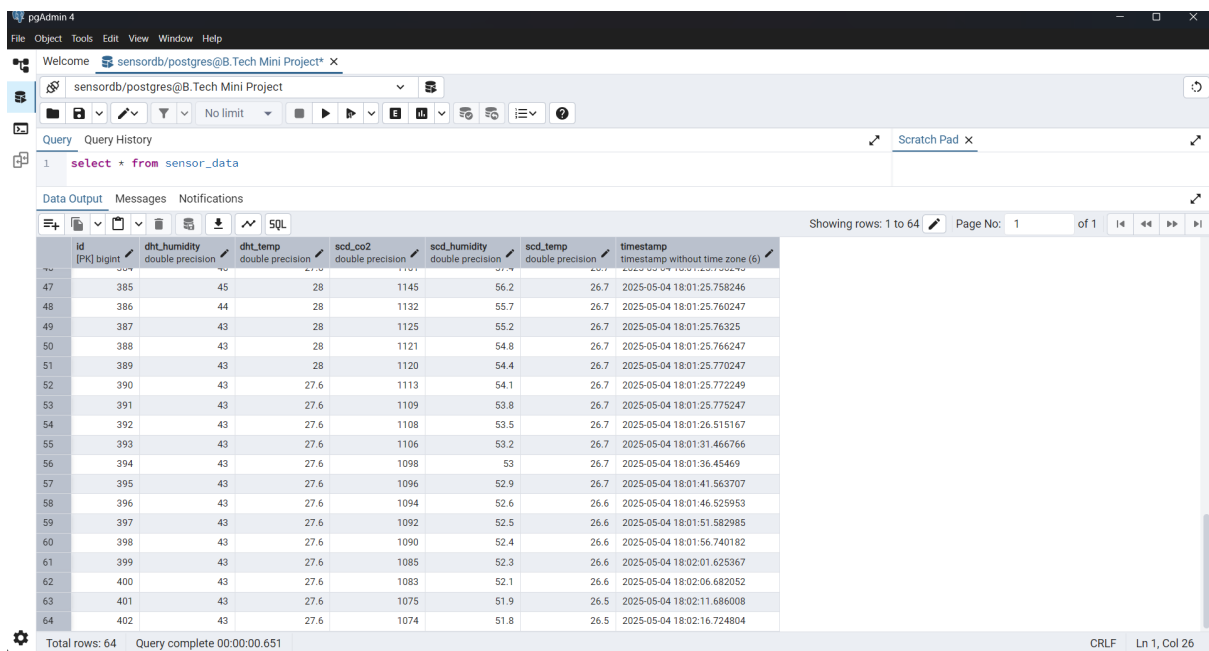
1. Run the SaveData application to store the data via API
2. Launch PGAdmin, connect to localhost:5432.
3. Create a database named `sensor_db`.
4. run the SQL Query

```
SELECT * from sensor_data
```

5. to delete the data fromt he database

```
DELETE from sensor_data
```

Spring Data Service This service receives sensor data via Kafka and stores it in PostgreSQL. Run the Spring Boot service on port 8082 to enable database storage.

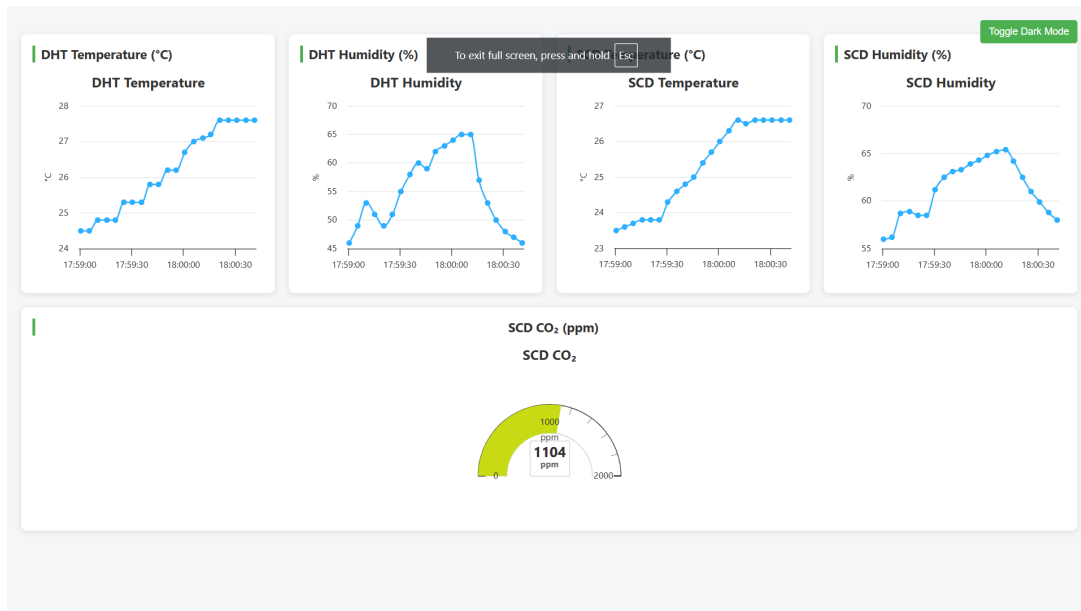


The screenshot shows the PGAdmin 4 interface. The top bar indicates the connection is to 'sensordb/postgres@B.Tech Mini Project'. The query editor shows the query 'select * from sensor_data'. The results pane displays a table with 64 rows. The columns are: id (PK), dht_humidity, dht_temp, scd_co2, scd_humidity, scd_temp, and timestamp. The data shows sensor readings from May 4, 2025, at 18:01:25.758246 to 18:02:16.724804.

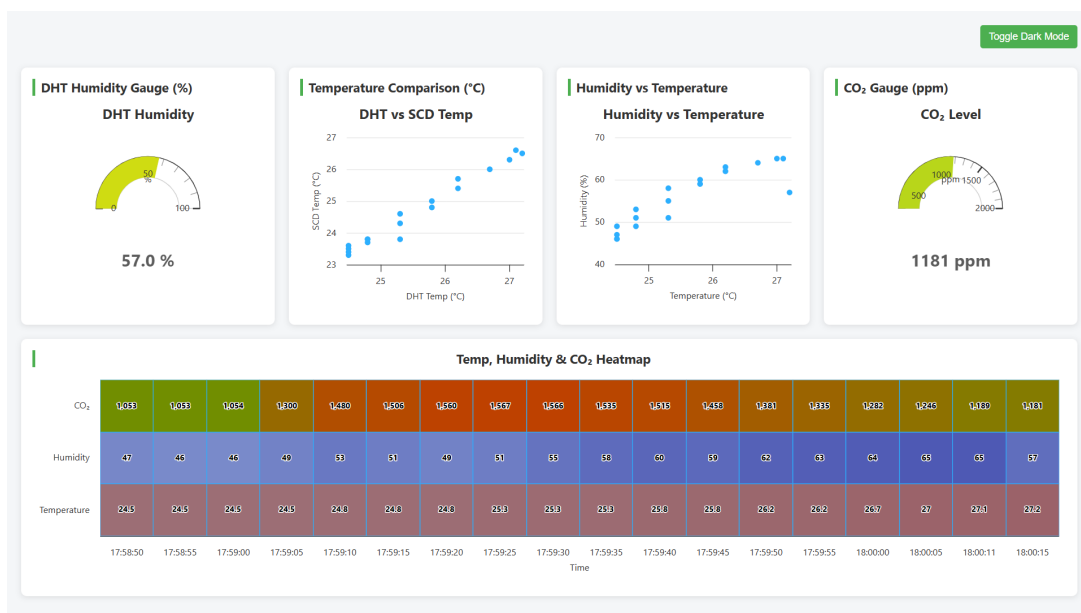
id (PK)	dht_humidity	dht_temp	scd_co2	scd_humidity	scd_temp	timestamp
385	45	28	1145	56.2	26.7	2025-05-04 18:01:25.758246
386	44	28	1132	55.7	26.7	2025-05-04 18:01:25.760247
387	43	28	1125	55.2	26.7	2025-05-04 18:01:25.76325
388	43	28	1121	54.8	26.7	2025-05-04 18:01:25.766247
389	43	28	1120	54.4	26.7	2025-05-04 18:01:25.770247
390	43	27.6	1113	54.1	26.7	2025-05-04 18:01:25.772249
391	43	27.6	1109	53.8	26.7	2025-05-04 18:01:25.775247
392	43	27.6	1108	53.5	26.7	2025-05-04 18:01:26.515167
393	43	27.6	1106	53.2	26.7	2025-05-04 18:01:31.466766
394	43	27.6	1098	53	26.7	2025-05-04 18:01:36.45469
395	43	27.6	1096	52.9	26.7	2025-05-04 18:01:41.563707
396	43	27.6	1094	52.6	26.6	2025-05-04 18:01:46.525953
397	43	27.6	1092	52.5	26.6	2025-05-04 18:01:51.582985
398	43	27.6	1090	52.4	26.6	2025-05-04 18:01:56.740182
399	43	27.6	1085	52.3	26.6	2025-05-04 18:02:01.625367
400	43	27.6	1083	52.1	26.6	2025-05-04 18:02:06.682052
401	43	27.6	1075	51.9	26.5	2025-05-04 18:02:11.686008
402	43	27.6	1074	51.8	26.5	2025-05-04 18:02:16.724804

5.2 Web Visualization

- **CO₂, Temperature, and Humidity Line Charts:** localhost:8081/basic-chart.html
 - Line charts for DHT11 Humidity & Temperature, SCD40 Humidity & Temperature.
 - A real-time gauge for CO₂ concentration levels.



- **Advanced Comparative & Distribution Visualization:** localhost:8081/advanced-chart.html
 - A scatter chart comparing Temperature vs. Humidity.
 - CO₂ and Humidity gauges.
 - A scatter chart comparing temperature between the SCD40 & DHT11 sensors.
 - A heatmap of the distribution of CO₂, Humidity, and Temp values over time.



5.3 Real-Time Smart Feedback

Threshold-based suggestions:

Parameter	Threshold	Action
CO ₂ Level	> 1200 ppm	Open windows or turn on exhaust fan
Temp	> 28 °C	Turn on fan or AC
Temp	> 20 °C	Use heater
Humidity	> 35%	Use humidifier
Humidity	> 65%	Use ventilation or dehumidifier

Table 2: Smart Feedback Rules

The screenshot shows an IDE with a project named 'RealTimeSensorDataStream'. The project structure includes a 'Consumer-3' module with a 'MakeDecision' sub-module. The 'MakeDecision' module contains a 'Service' interface and a 'KafkaConsumerService' implementation. The 'MakeDecisionApplication' class is the main entry point, which is a Spring Boot application. The code in the 'MakeDecisionApplication' class is as follows:

```

1 package kafka.consumer.MakeDecision;
2
3 import org.springframework.boot.SpringApplication;
4 import org.springframework.boot.autoconfigure.SpringBootApplication;
5 import org.springframework.boot.autoconfigure.jdbc.DataSourceAutoConfiguration;
6
7 @SpringBootApplication(exclude = {DataSourceAutoConfiguration.class})
8 public class MakeDecisionApplication {
9
10     public static void main(String[] args) { SpringApplication.run(MakeDecisionApplication.class, args); }
11
12 }
13
14
15

```

The 'Run' console shows the following logs:

```

2025-05-04T18:07:40.785+05:30 INFO 55516 --- [MakeDecision] [ntainer#0-0-C-1] o.a.k.c.c.internals.ConsumerUtils : Setting offset for partition mqtt_data_topic-2 to the
2025-05-04T18:07:40.785+05:30 INFO 55516 --- [MakeDecision] [ntainer#0-0-C-1] o.a.k.c.c.internals.ConsumerUtils : Setting offset for partition mqtt_data_topic-1 to the
2025-05-04T18:07:40.785+05:30 INFO 55516 --- [MakeDecision] [ntainer#0-0-C-1] o.a.k.c.c.internals.ConsumerUtils : Setting offset for partition mqtt_data_topic-0 to the
2025-05-04T18:07:40.786+05:30 INFO 55516 --- [MakeDecision] [ntainer#0-0-C-1] o.s.k.l.KafkaMessageListenerContainer : mqtt-data-group-3: partitions assigned: [mqtt_data_to

```

The logs also show several warnings: 'CO₂ high: Open windows or turn on exhaust fan'.

6 Additional Resources

- Kafka Downloads: <https://kafka.apache.org/downloads>
- Kafka Tutorial: https://www.youtube.com/watch?v=tU_37niRh4U

7 Troubleshooting

Symptom	Fix
No MQTT connection	Check ESP32 Wi-Fi HiveMQ credentials
No Kafka data	Ensure topic exists correct broker URL
Dashboard not loading	Verify Spring Boot service on port 8081
Firewall issues	Confirm port 9092 is open

Startup Sequence:

1. ESP32 publishes MQTT → Kafka
2. Start ZooKeeper
3. Start Kafka server
4. Run Kafka consumer console (test)
5. Launch Spring Boot apps (DB Dashboard)